

# CS10720 Problems and Solutions

Thomas Jansen

Today: Search in a Sorted Array:  
Linear and Binary Search

February 15<sup>th</sup>

# Organisational Issues

**Important** 'additional' lecture this Friday, 4-5pm, EL 0.26  
(replacing the Monday lecture that fell victim to the storm)

**Important** portfolio this week **same contents as always**  
lecture summary Monday and Thursday  
answer to practicals questions  
(**not** lecture summary Friday)  
**same deadline as always, Friday, 7pm**

**Remember** portfolio feedback for last week **later**  
(I'm working on it. It'll take more time.)

# Plans for Today

## ① Problem Definition

Motivation

Specifics

## ② Linear Search

Idea and Implementation

Assessment

## ③ Binary Search

Idea and Implementation

Assessment

## ④ Summary

Summary & Take Home Message

# Our Problem for Today

**General formulation** 'Find something you are looking for.'

**Problem** Have a large number of different items  
each with a unique identifier (called **key**)  
in some arrangement  
that is sorted in ascending order according to keys.  
You are given a key, you have to find the item.

**Observation** used to be frequent problem for people  
and still is frequent problem for computers

**Plan for today** learn how to solve  
one specific instance of this problem on a computer  
in several different ways



## Problem Definition

We simplify the problem by

- ① removing the items and only considering the keys
- ② having integers as keys

**Remark** simplifies writing down implementations of solutions  
 but makes the problem almost pointless  
 however it should be clear how to use different keys  
 and add back items

**Remember** we concentrate on the essence of the search problem  
 without being distracted by implementation details

**Input** array keys of type int of length  $n$   
 specific key of type int

**Output** if  $(\exists i \in \{0, 1, \dots, n-1\} : \text{keys}[i] = \text{key})$  then  $i$ , else  $-1$

```
#define n 100000 /* just an example size */
int keys[n]; /* the array of keys */
int key; /* the key we are looking for */
int search(int *keys, int size, int key);
```

## Towards a Solution

What do we want from a solution to this problem?

- ① It should be **correct**.      ② It should be **efficient**.

A few **stupid** suggestions

- Pick a random position  $i$ , if  $\text{keys}[i] == \text{key}$  then return  $i$   
else return  $-1$

**Observation** really **fast** but **with high probability not correct**  
(actually correct if key is not in keys; **but** otherwise  
correct only with probability  $1/n$ , thus **not correct** with  
probability  $1 - 1/n$ )

- Repeat  $n$  times: pick a random position  $i$ , if  $\text{keys}[i] == \text{key}$   
then return  $i$ . Return  $-1$  if key not found.

**Observation** up to  $n$  times slower and still **not correct**  
(if key present not correct with probability  
 $(1 - 1/n)^n \approx 1/e > 0.36$ )

- Check all positions in some order. Return position where key  
found or  $-1$  if not found.      **Observation correct**

## Implementing Our First Correct Idea

```
int search(int *keys, int size, int key) {  
    int i; /* index variable for a loop over keys */  
  
    for ( i=0; i<size; i++ ) {  
        if ( keys[i] == key )  
            return i;  
    }  
    return -1;  
}
```

Does this have to be C? Can't you explain this in English?

**Remark** most algorithms (including this one)  
in three versions in the **lecture notes**:  
English, pseudo-code, C code  
**Please**, read the lecture notes if something is unclear!

<http://onlinetted.com>

## Reconsidering Our First Correct Idea

```
int search(int *keys, int size, int key) {
    int i; /* index variable for a loop over keys */

    for ( i=0; i<size; i++ ) {
        if ( keys[i] == key )
            return i;
    }
    return -1;
}
```

**Observation** algorithm is **stupid!**

**Observation** once  $\text{keys}[i] > \text{key}$  we can stop looking  
because the keys only get bigger the further we get

**Consequence** stop as soon as possible, i. e. when  $\text{keys}[i] > \text{key}$



## Implementing Our First Correct Idea a Bit Better

```
int search(int *keys, int size, int key) {  
    int i; /* index variable for a loop over keys */  
  
    i=0; /* start searching at the beginning */  
    while ( (i<size) && (keys[i]<key) ) { /* search for key */  
        i++; /* move to next item */  
    }  
    if ( (i==size) || (keys[i]>key) ) /* 2 cases for key not found */  
        return -1;  
    else  
        return i;  
}
```

Observation

correct



## Alternative Implementation

```
int search(int *keys, int size, int key) {  
    int i; /* index variable for a loop over keys */  
  
    i=size-1; /* start searching at the end */  
    while ( (i>=0) && (keys[i]>key) ) { /* search for key */  
        i--; /* move to next item */  
    }  
    if ( (i==-1) || (keys[i]<key) ) /* 2 cases for key not found */  
        return -1;  
    else  
        return i;  
}
```

Observation

also correct ✓

## Comparing Two Implementations

```
int search(int *keys, int size, int key) {
    int i=0;
    while ( (i<size) && (keys[i]<key) )
        i++;
    if ( (i==size) || (keys[i]>key) )
        return -1;
    else
        return i;
}
```

```
int search(int *keys, int size, int key) {
    int i=size-1;
    while ( (i>=0) && (keys[i]>key) )
        i--;
    if ( (i==-1) || (keys[i]<key) )
        return -1;
    else
        return i;
}
```

Observation both correct ✓ Which one is more efficient?

### Observations

- In the best case both find the key in the first step
- In the worst case both search whole array and return -1
- In the average case key position (present or not present)

is at each position with equal probability

both look at  $\frac{1}{n} \cdot 1 + \frac{1}{n} \cdot 2 + \dots + \frac{1}{n} \cdot n = \frac{n(n+1)}{2n} = \frac{n+1}{2}$  keys  
( $n = \text{size}$ )

- in all cases both perform equal

## Comparing Two Implementations

```
int search(int *keys, int size, int key) {
    int i=0;
    while ( (i<size) && (keys[i]<key) )
        i++;
    if ( (i==size) || (keys[i]>key) )
        return -1;
    else
        return i;
}
```

```
int search(int *keys, int size, int key) {
    int i=size-1;
    while ( (i>=0) && (keys[i]>key) )
        i--;
    if ( (i===-1) || (keys[i]<key) )
        return -1;
    else
        return i;
}
```

### Observations

- both **correct**
- **In the best case** considering 1 key  
**very unlikely** and not very important
- **in the worst case** both considering  $n$  keys ( $n = \text{size}$ )
- **in the average case** both considering  $(n + 1)/2$  keys

Can we search significantly faster than this?

# Improving Over Linear Search

```
int search(int *keys, int size, int key) {
    int i=0;
    while ( (i<size) && (keys[i]<key) )
        i++;
    if ( (i==size) || (keys[i]>key) )
        return -1;
    else
        return i;
}
```

```
int search(int *keys, int size, int key) {
    int i=size-1;
    while ( (i>=0) && (keys[i]>key) )
        i--;
    if ( (i==-1) || (keys[i]<key) )
        return -1;
    else
        return i;
}
```

Do you search in a phonebook  
by reading it from front to back  
or by reading it from back to front?



**Observation** we can **exploit the order** of the keys  
much better than this  
by trying to locate the relevant area more quickly

## Binary Search

**Idea** try your luck in the middle and continue in the correct half using the same idea in that half and so on...

```
int search(int *keys, int size, int key) {
    int left, right; /* defining boundaries of search area */
    int middle; /* the 'middle' where we hope to find the key */
    left = 0; /* leftmost key */
    right = size-1; /* rightmost key */
    while ( left <= right ) {
        middle = left + (right-left)/2; /* midpoint */
        if ( keys[middle] == key )
            return middle;
        if ( keys[middle] < key ) /* need to look in right half */
            left = middle+1; /* can exclude middle */
        else /* need to look in left half */
            right = middle-1; /* can exclude middle */
    }
    return -1;
}
```

## Binary Search: Correctness

```
while ( left <= right ) {
    middle = left + (right-left)/2; /* midpoint */
    if ( keys[middle] == key )
        return middle;
    if ( keys[middle] < key ) /* need to look in right half */
        left = middle+1; /* can exclude middle */
    else /* need to look in left half */
        right = middle-1; /* can exclude middle */
}
return -1;
```

Is this correct? yes (but only 'hand waving' now; more in March)

- if key **not present**
  - code never returns anything but  $-1$
  - $(\text{right} - \text{left})$  gets smaller each round
  - when  $(\text{right} - \text{left}) < 0$  code returns  $-1$  ✓
- if key **present**
  - initially key is between left and right (inclusive)
  - this is always the case
  - at the latest when  $(\text{right} - \text{left}) = 0$  key is found and code returns correct position ✓

## Binary Search: Efficiency

```
int search(int *keys, int size, int key) {
    int left, right; /* defining boundaries of search area */
    int middle; /* the 'middle' where we hope to find the key */
    left = 0; /* leftmost key */
    right = size-1; /* rightmost key */
    while ( left <= right ) {
        middle = left + (right-left)/2; /* midpoint */
        if ( keys[middle] == key )
            return middle;
        if ( keys[middle] < key ) /* need to look in right half */
            left = middle+1; /* can exclude middle */
        else /* need to look in left half */
            right = middle-1; /* can exclude middle */
    }
    return -1;
}
```

Is this more efficient than linear search?

**Observation** in the **best case** considering 1 key  
(the same as for both variants of linear search)  
but equally **unlikely**

**Let's** only consider the **worst case**  
(since the average case is a bit more difficult)  
i. e., case where key is not present



## Binary Search: Efficiency in the Worst Case

```
int search(int *keys, int size, int key) {
    int left, right; /* defining boundaries of search area */
    int middle; /* the 'middle' where we hope to find the key */
    left = 0; /* leftmost key */
    right = size-1; /* rightmost key */
    while ( left <= right ) {
        middle = left + (right-left)/2; /* midpoint */
        if ( keys[middle] == key )
            return middle;
        if ( keys[middle] < key ) /* need to look in right half */
            left = middle+1; /* can exclude middle */
        else /* need to look in left half */
            right = middle-1; /* can exclude middle */
    }
    return -1;
}
```

How many keys are considered when the key is not present?

- in each round 2 keys (the same key twice)
- number of keys is twice the number of rounds
- What is the number of rounds in the worst case?
- Remember either left changes to  $\text{middle} + 1$  or right changes to  $\text{middle} - 1$  with  $\text{middle} = \text{left} + (\text{right} - \text{left})/2$

# Binary Search: Number of Rounds in the Worst Case

**Remember**    **either** left changes to  $\text{middle} + 1$   
                   **or** right changes to  $\text{middle} - 1$   
                   with  $\text{middle} = \text{left} + (\text{right} - \text{left})/2$

**In the following**     $l = \text{left}, r = \text{right}, n = \text{size}$

**Case 1**    range changes from  $r - l$  to  $r - \left(l + \left\lfloor \frac{r-l}{2} \right\rfloor + 1\right)$   
 $= (r - l) - \left(\left\lfloor \frac{r-l}{2} \right\rfloor + 1\right) = (r - l) - \left\lfloor \frac{r-l}{2} \right\rfloor - 1$   
 $\leq (r - l) - \frac{r-l}{2} = \frac{r-l}{2}$

**Case 2**    range changes from  $r - l$  to  $\left(l + \left\lfloor \frac{r-l}{2} \right\rfloor - 1\right) - l$   
 $= \left\lfloor \frac{r-l}{2} \right\rfloor - 1 < \frac{r-l}{2}$

**Remember**    in each round range shrinks from  $r - l$  to  $\leq (r - l)/2$   
                   and round with  $r - l < 1$  is last (in the worst case)

# Computing Number of Rounds in the Worst Case

**Remember** in each round range shrinks from  $r - l$  to  $\leq (r - l)/2$   
 and round with  $r - l < 1$  is last (in the worst case)  
 initially  $r - l = n - 1$  (Assuming  $r - l = n$  makes it worse)

**Consider** development of range  $r - l$  (in the worst case)

$$n \rightsquigarrow \frac{n}{2} \rightsquigarrow \frac{n}{4} \rightsquigarrow \frac{n}{8} \rightsquigarrow \frac{n}{16} \rightsquigarrow \dots \rightsquigarrow X$$

$$n \rightsquigarrow \frac{n}{2^1} \rightsquigarrow \frac{n}{2^2} \rightsquigarrow \frac{n}{2^3} \rightsquigarrow \frac{n}{2^4} \rightsquigarrow \dots \rightsquigarrow \frac{n}{2^k}$$

$$\text{with } X < 1 \quad \text{with } 2^k > n$$

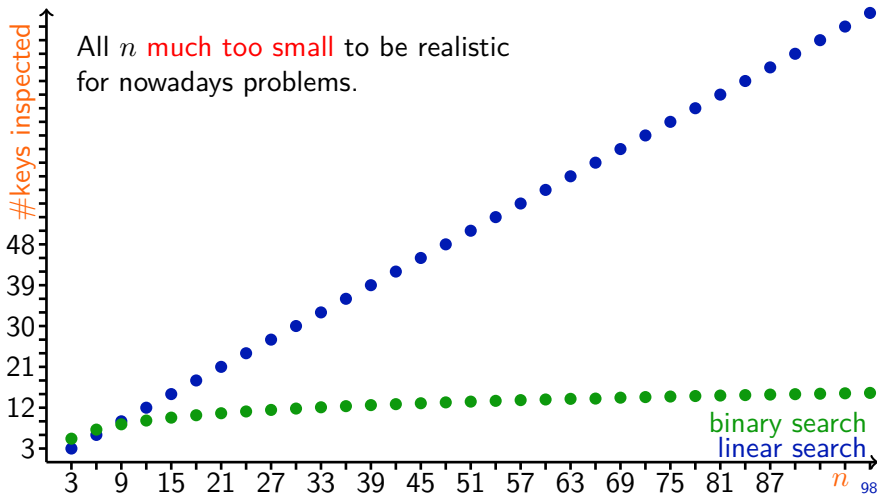
so  $k = 1 + \log_2 n$  rounds are enough in the worst case

total number of rounds  $\leq 1 + \log_2 n$  in the worst case

**Summary** binary search considers at most  $2\log_2(n) + 2$  keys  
 while linear search considers up to  $n$  keys

# Comparing Linear and Binary Search

**Summary**    binary search considers at most  $2\log_2(n) + 2$  keys  
while linear search considers up to  $n$  keys



# Comparing Linear and Binary Search for Larger Arrays

**Summary**    binary search considers at most  $2\log_2(n) + 2$  keys  
while linear search considers up to  $n$  keys

**Comparison**    for larger  $n$

$n$	linear	binary
10	10	< 9
100	100	< 16
1 000	1 000	< 22
10 000	10 000	< 29
100 000	100 000	< 36
1 000 000	1 000 000	< 42
10 000 000	10 000 000	< 49
100 000 000	100 000 000	< 56
1 000 000 000	1 000 000 000	< 62

# Summary & Take Home Message

## Things to remember

- searching in a **sorted** array
- linear search
- linear search with linear run time in the worst and average case
- binary search
- binary search with logarithmic time in the worst case

## Take Home Message

- It pays not to be very stupid.
- It pays to exploit structural properties.
- Binary search is very, very, very much faster than linear search unless we search in very small arrays.

**Lecture feedback** <http://onlinetted.com>